

# Spark Streaming



<http://spark.apache.org/docs/latest/streaming-programming-guide.html>

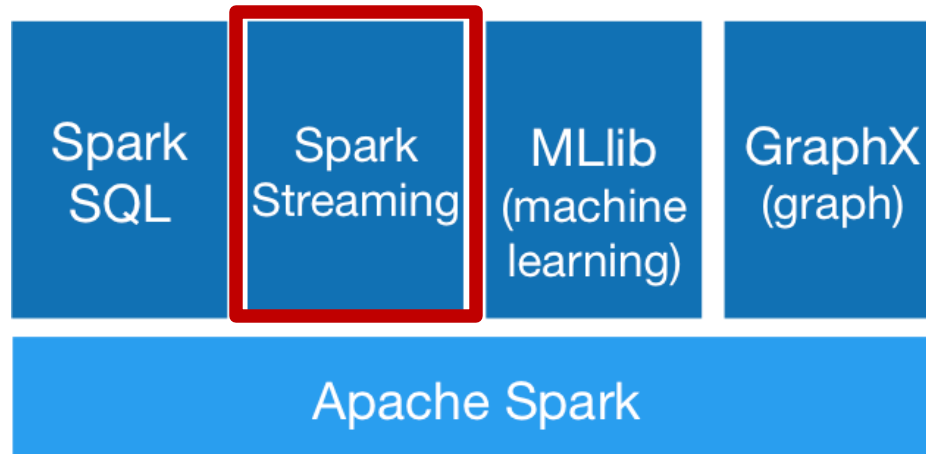
**Nagiza F. Samatova**, [samatova@csc.ncsu.edu](mailto:samatova@csc.ncsu.edu)

Professor, Department of Computer Science  
North Carolina State University

Senior Scientist, Computer Science & Mathematics Division  
Oak Ridge National Laboratory

# Spark Streaming

- Streaming applications benefit from acting on data as soon as it arrives by:
  - tracking statistics about page views in real time,
  - training a machine learning model, or
  - automatically detecting anomalies
- **Spark Streaming** is Spark's module for such applications:
  - lets users write streaming applications using a very similar API to batch jobs, and
  - thus reuse a lot of the skills and even code they built for those.
- **DStreams**: Spark's abstraction for discretized streams



# DStreams

- **DStream** is a sequence of **RDDs** arriving at each time step (hence, the name “discretized”).
- DStreams can be created from various input sources, such as Flume, Kafka, or HDFS.
- Operations on DStreams:
  - **Transformations**: yield a new DStream
  - **Actions**: operate on RDDs of the DStreams
  - **Sliding Window** operations
  - **Output operations**: write data to an external system
- **24/7 Operation**: Unlike batch programs, Spark Streaming applications operate 24/7:
  - via **Checkpointing** mechanism for storing data in a reliable file system such as HDFS
  - via **Restarting** (manually or automatically) applications on failure

# Example: Step 1: Import Modules

---

**Example:** Receive a stream of newline-delimited lines of text from a server running at port 77777, filter only the lines that contain the word “error,” and print them

# Step 1. Import the required modules

```
from __future__ import division
from pyspark import SparkConf, SparkContext
from pyspark.streaming import StreamingContext
from pyspark.streaming.kafka import KafkaUtils
import sys
import random
```

# Step 2: Create StreamingContext

- Create a **StreamingContext**, key entry point for streaming functionality
  - set up an underlying SparkContext to process the data
- StreamingContext takes as input a **batch interval** specifying how often to process new data (e.g., every second)

# Step 2. Create StreamingContext

```
conf = (SparkConf()  
        .setMaster("local[4]")  
        .setAppName("SentimentAnalysis")  
        .set("spark.executor.memory", "2g"))  
  
sc = SparkContext(conf=conf)  
  
ssc = StreamingContext(sc, 1)    # Create a streaming context with  
batch interval of 1 sec  
ssc.checkpoint("checkpoint")
```

# Step 3: Create DStream w/ socketTextStream

- Use `socketTextStream()` to create **DStream** based on text data received on port 7777 of the local machine
- Transform the DStream with `filter()` to get only lines that contain "error"
- Apply the output operation `pprint()` [pretty print] to print some of the filtered lines.

# Step 3. Create DStream

```
# Streaming filter for printing lines containing "error" in Python
lines = ssc.socketTextStream("localhost", 7777)
error_lines = lines.filter(lambda x: "error" in x)
error_lines.pprint()

# Start our streaming context and wait for it to "finish"
ssc.start()

# Wait for the job to finish
ssc.awaitTermination()
```

# Step 4: Test

- Start the Spark job, provide its input via a socket using **ncat**:
- Install: `# sudo apt-get install nmap`
- Start the socket on Ubuntu: `$ nc -lk 7777`
- Type in the lines, the lines that contain "error" will be printed within a respective interval on the terminal where this Spark program is running

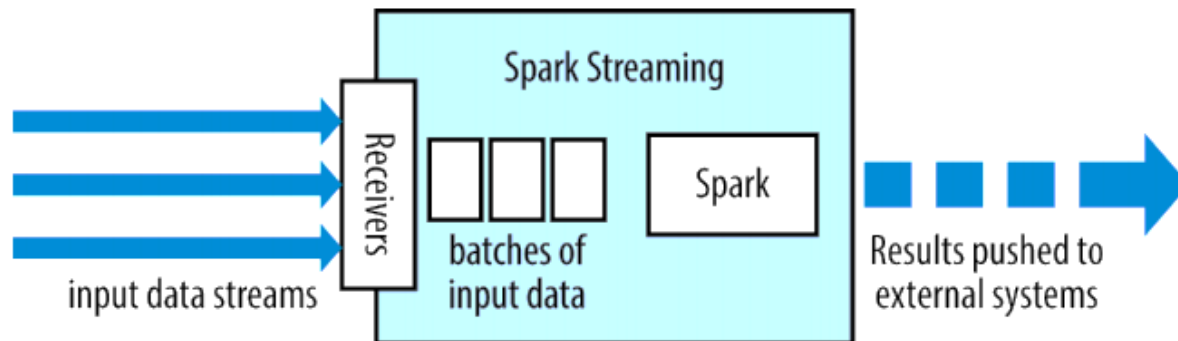
```
tolik@tolik-VirtualBox:~$ nc -lk 7777
error
another error
```

```
16/01/04 01:00:35 INFO DAGScheduler: Job 1 finished: runJob at PythonRDD.scala:93, took 5.084788 s
-----
Time: 2016-01-04 01:00:30
-----
another error

16/01/04 01:00:35 INFO JobScheduler: Finished job streaming job 1451887230000.0 from job set of time 1451887230000 ms
```

# Micro-Batching Architecture


- Spark Streaming uses a “**micro-batch**” architecture:
  - streaming computation is treated as a continuous series of batch computations on small batches of data
- Spark Streaming receives data from various input sources and groups it into small batches
  - New batches are created at regular time intervals
  - At the beginning of each time interval a new batch is created, and any data that arrives during that interval gets added to that batch
  - At the end of the time interval the batch is done growing
  - The size of the time intervals is called the **batch interval**:
    - The batch interval is typically between **500 milliseconds and several seconds**, as configured by application developer





# DStream as a Continuous Series of RDDs

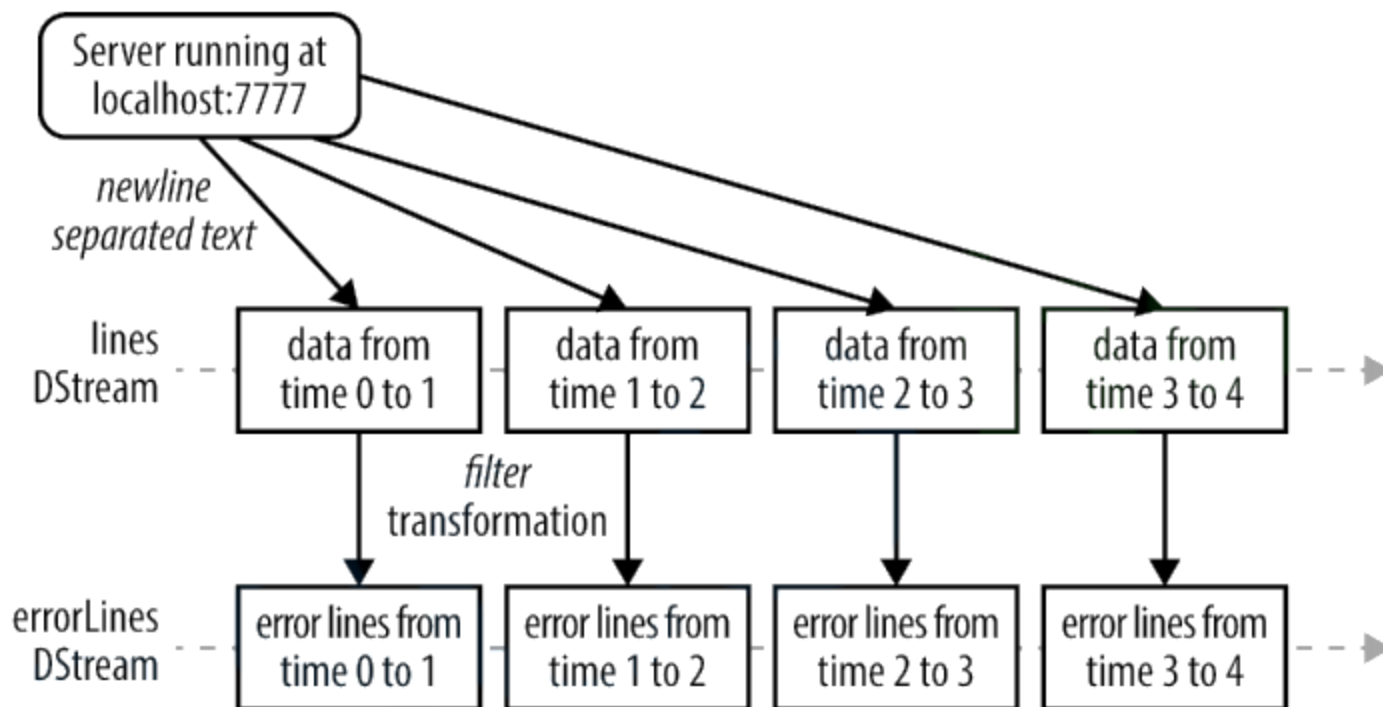
- DStream is a discretized stream, the programming abstraction in Spark Streaming
- DStream is a sequence of RDDs, where each RDD has one time slice of the data in the stream



- **Create DStreams:**
  - either from external input sources, or
  - by applying transformations to other DStreams
- **DStreams support many of the transformations on RDDs**
- **DStreams have *stateful* transformations to aggregate data over time**

# DStreams and Transformations Example

- In our example, we created a DStream from data received through a socket, and then applied a filter() transformation to it.
- This internally creates RDDs:




# Output Operations

---

- DStreams support output operations such as `pprint()`
- Output operations are similar to RDD actions in that they write data to an external system, but in Spark Streaming they **run periodically on each time step**, producing **output in batches**

# Execution of Spark Streaming

- The execution of Spark Streaming within Spark's driver-worker components



- For each input source, Spark Streaming launches **receivers**, which are tasks running within the application's executors that collect data from the input source and save it as RDDs

# Fault Tolerance in Spark Streaming

- The input data is **replicated** (by default) to another executor for **fault tolerance**:
  - data is stored **in-memory of the executors** as cached RDDs.
- StreamingContext in the driver program:
  - periodically runs Spark jobs to process this data and
  - combines it with RDDs from previous time steps
- Spark Streaming fault-tolerance for DStreams:
  - as long as a copy of the input data is available, it re-computes any state derived from it using the **lineage** of the RDDs
    - i.e. by rerunning the operations used to process it
- Spark Streaming also includes a mechanism called **checkpointing** that **saves state** periodically **to a reliable filesystem** (HDFS or S3)
  - when recovering lost data Spark Streaming needs only to go to the last checkpoint (typically 5-10 batches)

# Transformations

- Transformations on DStreams can be grouped into either **stateless** or **stateful**:
  - In **stateless transformations** the processing of **each batch** does not depend on the data of its previous batches
    - E.g., `map()`, `filter()`, and `reduceByKey()`
  - Although they look operating on the whole stream, internally each DStream is composed of multiple RDDs (batches), and each stateless transformation is applied separately **to each RDD**.
    - E.g., `reduceByKey()` will reduce data within each time step, but not across time steps.
  - **Stateful transformations** use data or intermediate results from previous batches to compute the results of the current batch
    - Transformations based on **sliding windows** and **tracking state** across time

# Stateless Transformations

Function name	Purpose	Scala example	Signature of user-supplied function on DStream[T]
map()	Apply a function to each element in the DStream and return a DStream of the result.	<code>ds.map(x =&gt; x + 1)</code>	<code>f: (T) → U</code>
flatMap()	Apply a function to each element in the DStream and return a DStream of the contents of the iterators returned.	<code>ds.flatMap(x =&gt; x.split(" "))</code>	<code>f: T → Iterable[U]</code>
filter()	Return a DStream consisting of only elements that pass the condition passed to filter.	<code>ds.filter(x =&gt; x != 1)</code>	<code>f: T → Boolean</code>
repartition()	Change the number of partitions of the DStream.	<code>ds.repartition(10)</code>	N/A
reduceByKey()	Combine values with the same key in each batch.	<code>ds.reduceByKey(   (x, y) =&gt; x + y)</code>	<code>f: T, T → T</code>
groupByKey()	Group values with the same key in each batch.	<code>ds.groupByKey()</code>	N/A

# Stateless Example

Example: log processing program that uses `map()` and `reduceByKey()` to **count log events by IP address** in each time step

```
# map() and reduceByKey() on DStream
# Run this example, and then copy the file to
directory = sys.argv[1]
print(directory)

# create DStream from text file
# Note: the spark streaming checks for any updates to this directory.
# So first, start this program, and then copy the log
# file logs/access_log.log to 'directory' location
log_data = ssc.textFileStream(directory)

# Parse each line using a utility class
access_log_dstream = log_data.map(ApacheAccessLog.parse_from_log_line)
    .filter(lambda parsed_line: parsed_line is not None)

# map each ip with value 1. So the stream becomes (ip, 1)
ip_dstream = access_log_dstream.map(lambda parsed_line: (parsed_line.ip, 1))

ip_count = ip_dstream.reduceByKey(lambda x,y: x+y)
ip_count.pprint(num = 30)
```



# Stateless Transformation over Multiple DStreams

---

- Stateless transformations can also combine data from multiple DStreams within each time step
  - For example, key/value DStreams have the same join-related transformations as RDDs
    - `cogroup()`, `join()`, `leftOuterJoin()`

# Join between two DStreams Example

- Given data keyed by IP address, join the request count against the bytes transferred

```
# Join two Dstreams

ip_bytes_dstream = access_log_dstream.map(lambda parsed_line:
    (parsed_line.ip, parsed_line.content_size))
ip_bytes_sum_dstream = ip_bytes_dstream.reduceByKey(lambda x,y: x+y)
ip_bytes_request_count_dstream = ip_count.join(ip_bytes_sum_dstream)
ip_bytes_request_count_dstream.pprint(num = 30)
```

- Or merge the contents of two different DStreams using the `union()` operator as in regular Spark, or using `StreamingContext.union()` for multiple streams

# transform()

- **transform():** DStreams provide an advanced operator that lets you operate directly on the RDDs inside them:
  - lets you provide any arbitrary RDD-to-RDD function to act on the DStream
- **Example application:**
  - to reuse batch processing code written for RDDs
  - For example, extractOutliers() function that acted on an RDD of log lines to produce an RDD of outliers (perhaps after running some statistics on the messages) can be reused within a transform()

```
def extractOutliers(rdd):  
    """ Currently, no logic implemented. But you can specify any rdd logic here """  
    return rdd  
  
transformed_access_log_dstream = access_log_dstream.transform(extractOutliers)  
transformed_access_log_dstream.pprint()
```

# Stateful Transformations

- **Stateful transformations** are operations on DStreams that track data across time
  - some data from previous batches is used to generate the results for a new batch
- The two main types are:
  - **windowed operations**: act over sliding window of time periods
  - **updateStateByKey()**: track state across events for each key
    - e.g. to build up an object representing each user session
- Stateful transformations require **checkpointing** to be enabled in your StreamingContext for fault tolerance

```
ssc.checkpoint("hdfs://...")
```


# Windowed Transformations

---

- **Windowed operations** compute the results in a longer time period than the StreamingContext's batch interval, by **combining results from multiple batches**.
  - Used to keep track of the most common response codes, content sizes, and clients in a web server access log
  - Need two parameters, **window duration** and **sliding duration**, each is a multiple of the StreamingContext's batch interval

# Windowing Parameters

- The **window duration** controls how many previous batches are considered, namely `windowDuration/batchInterval`
  - Eg: For a DStream with a batch interval of 10 seconds and a sliding window of the last 30 seconds (or last 3 batches), set the `windowDuration` to 30 seconds
- The **sliding duration**, defaults to the batch interval, controls how frequently the new DStream computes results.
  - For a DStream with a batch interval of 10 seconds, to compute out window only on every second batch, set the sliding interval to 20 seconds



every two time steps,  
compute a result over the  
3 previous time steps

# Request a window with window()

- **window()**: returns a new DStream with the data for the requested window
  - each RDD in the DStream resulting from window() will contain data from multiple batches that can be processed with count(), transform(), and so on.

```
# How to use window() to count data over a window
access_logs_window = access_log_dstream.window(windowDuration = 6,
        slideDuration=4)
window_counts = access_logs_window.count()
print( " Window count: ")
window_counts.pprint()
```

# Aggregate with `reduceByWindow()`



---

- **`reduceByWindow()`** and **`reduceByKeyAndWindow()`** allow to perform each window more efficiently
  - take a single reduce function (e.g., `+`) to run on the whole window
  - allow Spark to compute the reduction incrementally, by considering only data coming into the window.



# Naïve vs Incremental reduceByWindow()

- Difference between naïve reduceByWindow() and incremental reduceByWindow(), using an inverse function



# Example: reduceByWindow()

**Example: Use reduceByWindow() and reduceByKeyAndWindow() to count visits by each IP address more efficiently**

```
# Ip counts per window
ip_count_dstream = ip_dstream.reduceByKeyAndWindow(func = lambda x,y:
    x+y, invFunc = lambda x,y: x-y, windowDuration = 6, slideDuration=4)
ip_count_dstream.pprint(num=30)
```

# Counting Data in DStreams

- **countByWindow()**: gives a DStream representing the number of elements in each window
- **countByValueAndWindow()**: gives a DStream with the counts for each value

```
# Windowed count operation
ip_dstream = access_log_dstream.map(lambda entry: entry.ip)
ip_address_request_count =
    ip_dstream.countByValueAndWindow(windowDuration = 6, slideDuration=4)
ip_address_request_count.pprint()
request_count =
    access_log_dstream.countByWindow(windowDuration = 6, slideDuration=4)
request_count.pprint()
```

# updateStateByKey()

- **updateStateByKey()** enables to maintain state across the batches in a DStream by providing **access to a state variable** for DStreams of key/value pairs:
  - e.g. to track sessions as users visit a site
- Given a DStream of (key, event) pairs, it lets you construct a new DStream of (**key, state**) pairs by **taking a function** that specifies **how to update the state for each key given new events**:
  - e.g. in a web server log, events might be visits to the site, where the key is the user ID
  - Using updateStateByKey() we could track the last 10 pages each user visited. This list would be the “state” object, and will be updated as each event arrives

# How to Use `updateStateByKey()`?

- Provide a function `update(events, oldState)` that takes in the events that have arrived for a key and its previous state, and returns a `newState` to store for it.
  - `events` is a list of events that arrived in the current batch (may be empty)
  - `oldState` is an optional state object, stored within an `Option`; it might be missing if there was no previous state for the key
  - `newState` returned by the function is also an `Option`; we can return an empty `Option` to specify that we want to delete the state
- The result of `updateStateByKey()` will be a new `DStream` that contains an **RDD of (key, state) pairs** on each time step

# Example: updateStateByKey()

- **updateStateByKey()** is used to keep a running count of the number of log messages with each HTTP response code.
  - keys are the response codes
  - state is an integer representing each count
  - events are page views
- Unlike window examples it keeps “infinitely growing” count since the beginning of the program

```
# Running count of response codes using updateStateByKey() in Scala
# This basically maintains a running sum , rather than sum in windows

def state_full_sum(new_values, global_sum):
    return sum(new_values) + (global_sum or 0)

response_code_dstream = access_log_dstream.map(lambda entry: (entry.response_code, 1))
response_code_count_dstream = response_code_dstream.updateStateByKey(state_full_sum)
response_code_count_dstream.pprint()
```

# Output Operations

---

- Output operations specify what needs to be done with the final transformed data in a stream
  - e.g. pushing it to an external database or printing it to the screen
- Much like lazy evaluation in RDDs, **if no output operation is applied on a DStream and any of its descendants, then those DStreams will not be evaluated.**
- If there are no output operations set in a StreamingContext, then the context will not start

# Output Example

---

- Spark Streaming has similar to `save()` operations for DStreams, each of which takes a directory to save files into and an optional suffix.
- The results of **each batch** are saved as **subdirectories** in the given directory, with the time and the suffix in the filename

```
ip_address_request_count.saveAsTextFiles(prefix = "outputDir", suffix = "txt")
```



# Acknowledgement

---

- Nirmesh Khandelwal, NCSU
- Anatoli Melechko, NCSU
- Learning Spark by Holden Karau, Andy Konwinski, Patrick Wendell, Matei Zaharia  
(<http://shop.oreilly.com/product/0636920028512.do>)