


# Time Series (TS) Forecasting



*How meaningful is your forecasting?*

**Nagiza F. Samatova**, [samatova@csc.ncsu.edu](mailto:samatova@csc.ncsu.edu)

Professor, Department of Computer Science  
North Carolina State University

Senior Scientist, Computer Science & Mathematics Division  
Oak Ridge National Laboratory

# Learning Objectives

---

- Introduce the principles of forecasting
- Learn how to use forecasting effectively
- Learn when to use forecasting

# Recommended Resources

---

- **Books**

- Free and online ([otexts.com/fpp](http://otexts.com/fpp)): Forecasting Principles & Practice by R. Hyndman, G. Athanasopoulos ← **Excellent Book!!!**
- Practical Time Series Forecasting with R: A Hand-on Guide by Shmueli & Lichtendahl

- **Packages**

- R: `fpp` (`install.packages("fpp", dependencies=TRUE)`)s

# Get Familiar with the Package

---

```
1 install.packages("fpp", dependencies=TRUE)
2 library(fpp)
3
4 help.search("forecasting")
5 help(forecast)
6 example("forecast.ar")
7
8 # similar names
9 apropos("forecast")
10
11 help(package="fpp")
```


# Dependency Packages for "fpp"

---

- **library (fpp): will load the following**
  - data sets and examples
  - **forecast** package: forecasting functions
  - **tseries** packages: some time series functions
  - **fma** package: lots of time series data
  - **expsmooth** package: ts data
  - **lmtest** package: for some regression functions
- **data (package="fpp"): show available ts data**
- **data (package="fma"): show available ts data**

# Use Case: Airline Passenger Traffic Forecast

- Problem: Forecast passenger traffic on major airlines



- Large amount of data on previous routes is available.
- Traffic is affected by school holidays, special events, advertising campaigns, competition, etc.

# Other Use Cases for TS Forecasting

---

- Daily Stock Prices (e.g., Dow Jones Index)
  - Monthly, quarterly and annual profits
  - Monthly, quarterly and annual product demands
  - Quarterly beer production
  - Monthly rainfall
  - Monthly residential electricity sales
- 
- **library (fpp): will load the following**
  - **data (package="fpp"): show available ts use cases**
  - **data (package="fma"): show available ts use cases**

# Time Series Forecasting: Problem

- Time series:
  - A sequence(s) of observations collected over time.
- Assumptions:
  - The time periods are **equally spaced** (e.g., not always true: Bitcoin data was changed to fit this assumption).

**Forecasting** is estimating how the sequence of observations will continue into the future.



# Basic Notation

Symbol	Definition
$t = 1, 2, 3, \dots,$	An index for the time period of interest; e.g., for a <i>daily</i> time period, $t = 1$ means day 1, $t = 2$ means day 2, etc.
$y_1, y_2, \dots, y_T$	A series of $T$ values measure over $T$ time periods; e.g., for the annual average stock price, $y_1$ denotes the price for year 1, $y_2$ denotes the price for year 2, etc.
$F_t$ or $\hat{y}_t$	The forecast value for time period $t$
$F_{t+k}$ or $\widehat{y_{t+k}}$	The $k$ -step-ahead forecast when forecasting time is $t$ ; e.g., $F_{t+1}$ is the forecast for time period $(t + 1)$ made during the time period $t$
$e_t = y_t - F_t$	The forecast error for time period $t$

# Baseline: Simple Forecasting Methods

- **Average:** `meanf ( ts.data, h=20 )`
  - Forecast of all future values is the mean of historical data  $\{y_1, \dots, y_T\}$
  - $F_{T+h} = \hat{y}_{T+h} = \bar{y} = (y_1 + \dots + y_T)/T$
- **Naive:** `naive (ts.data, h=20)` or `rwf (ts.data, h=20)`
  - Forecast is equal to the last observed value
  - $F_{T+h|T} = \hat{y}_{T+h|T} = y_T$
- **Seasonal naive:** `snaive (ts.data, h=20)`
  - Forecast is equal to the last value from the same season
  - $\hat{y}_{T+h|T} = y_{T+h-km}$ , where  $m$  is the seasonal period and  $k = \text{round}\left(\frac{h-1}{m}\right) + 1$
- **Drift:** `rwf (ts.data, drift=TRUE, h=20)`
  - Forecast is equal to the last value plus the average change
  - Equivalent to extrapolating a line between the first and last observation
  - $F_{T+h|T} = \hat{y}_{T+h|T} = y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1}) = y_T + \frac{h}{T-1} (y_T - y_1)$

# Ex: Show Baseline Forecasts for TS Data

---

- `library(fpp)`
- `data (package="fma")`
- `data (package="fpp")`
  
- `ts.data <- data (beer)`

# Visual Analysis of Time Series (TS) Data

- **Time plots:** `plot` or `plot.ts` (e.g., `plot(a10)`)
- **Seasonal plots:** `seasonplot` (e.g., `seasonplot(a10)`)
  - Data from each season is overlapped.
  - To view the underlying seasonal patterns.
- **Seasonal subseries plots:** `monthplot` (e.g., `monthplot(a10)`)
  - Data for each season collected together in time plot as separate time series.
  - To view the underlying seasonal patterns and the changes in seasonality over time
- **Lag plots:** `lag.plot`
- **ACF plots:** `Acf`

```
beer <- window(ausbeer, start=1992)

plot(beer)

seasonplot(beer, year.labels=TRUE)

monthplot(beer)
```

# Ex: Visually Explore Different TS Data Sets

---

- `library(fpp)`
- `data (package="fma")`
- `data (package="fpp")`
  
- `ts.data <- data (beer)`
- `***plot***(ts.data)`

# TS Parts: **Systematic** vs **Non-systematic**

TS Part	Definition	Detection	How to deal w/
<b>Level</b>	Average value of ts		
<b>Trend</b>	Long-term increase decrease in the data	lag.plot	De-trend via lag-1 differencing
<b>Seasonality</b>	Variations occurring during known periods of the year (monthly, quarterly, holidays)	lag.plot, Acf plots	De-seasonalize via lag-k differencing
<b>Cycles</b>	Other oscillating patterns about the trend (e.g., business or economic conditions)		
<b>Auto-correlation</b>	Correlation between neighboring points in ts	Acf, lag.plot	
<b>Noise</b>	Residuals after level, trend, seasonality, and cycles are removed	Normality tests	

# Non-systematic Part: Noise, or Residuals

**Residuals:** difference between observed value & its forecast based on all previous observations:  $e_t = y_t - \hat{y}_{t|t-1}$

- Assumptions

- $\{e_t\}$  are **uncorrelated**: otherwise, information is left in residuals that should be used in computing forecasts
- $\{e_t\}$  have **zero means**; otherwise, forecasts are biased
- $\{e_t\}$  have **constant variance**
- $\{e_t\}$  are **normally distributed**


# Seasonal vs. Cyclic Patterns

- Seasonal pattern exists when a time-series is influence by seasonal factors
  - quarter of the year, month, day of the week, holidays
- Cyclic patterns exist when data exhibit rises & falls that are not of fixed period
  - usually of at least 2 years


Pattern	Seasonal	Cyclic
<b>Length</b>	Constant length	Variable length
<b>Average length</b>	Shorter	Longer
<b>Magnitude</b>	Magnitude of seasonal pattern is less variable	More variable
<b>Predictability of peaks and troughs</b>	Predictable	Hard to predict for long-term cyclic data




# Cyclic: Sales of new one-family houses in US




# Seasonal: Annual Lynx Trappings in Canada




# Trend & Seasonality: Electricity Production



# Trend Pattern: US Treasury Bill Contracts



# Upward Trend: Australian GDP



# Additive and Multiplicative TS Components

---

- A time series with **additive** components can be modeled as:

$$y_t = \textit{Level} + \textit{Trend} + \textit{Seasonality/Cycles} + \textit{Noise}$$

- A time series with **multiplicative** components is modeled as:


$$y_t = \textit{Level} \times \textit{Trend} \times \textit{Seasonality/Cycles} \times \textit{Noise}$$

- Forecasting methods attempt to isolate the systematic part and quantify the noise level.
  - The systematic part is used for generating point forecasts
  - The level of noise helps assess the uncertainty associated with the point forecasts


# Trend & Seasonality Patterns

- Trend patterns are commonly approximated by linear, exponential and other mathematical functions


Constant Trend  
Non-seasonal




Constant Trend with  
Multiplicative Seasonality




Upward Linear Trend  
Non-seasonal




Upward Linear Trend with  
Additive Seasonality




Constant Trend with  
Additive Seasonality




Upward Exponential Trend  
Non-seasonal



Upward Linear Trend with  
Multiplicative Seasonality



Upward Exponential Trend with  
Additive Seasonality



# Autocorrelation: Motivation

---

- When we use linear regression for time series forecasting, we are able to account for patterns such as trend and seasonality.
- However, ordinary regression models do not account for dependence between observations, which in cross-sectional data is assumed to be absent.
- Yet, in the time series context, observations in neighboring periods tend to be correlated.
- Such correlation or autocorrelation, is informative and can help in improving forecasts.
- How do you compute autocorrelation and how do you best utilize the information for improving forecasting?



# Autocorrelation

**Autocorrelation:** measures linear relationship between lagged values of the time series

- Lag-1:  $y_t$  and  $y_{t-1}$
- Lag-2:  $y_t$  and  $y_{t-2}$
- Lag-6:  $y_t$  and  $y_{t-6}$


# Computing Autocorrelation

---

- Correlation between values of a time series in neighboring periods is called autocorrelation because it describes a relationship between the series and itself.
- To compute autocorrelation, we compute the correlation between the series and a lagged version of the series.
  - A lagged series is a copy of the original series which is moved forward one or more time periods.
  - A lagged series with lag-1 is the original series moved forward one time period. A lagged series with lag-2 is the original series moved forward two time periods, and so on.
- The lag-1 autocorrelation is the correlation
  - measures the linear relationship between values in consecutive time periods
  - is computed correlation between the original series and the lag-1 series

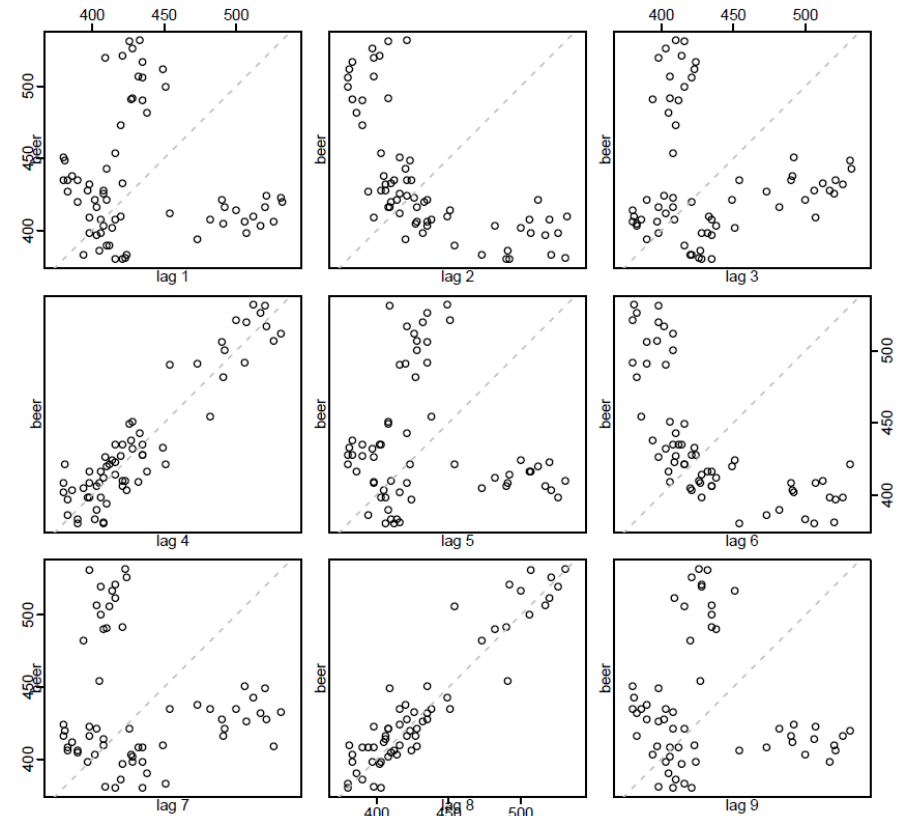
# Linear Relationship between Lagged Values

```
> lag.plot(beer, lags=9)
```



# Lagged Scatterplots

```
> lag.plot(beer, lags=9, do.lines=FALSE)
```



- Each graph shows  $y_t$  plotted against  $y_{t-k}$  for different values of  $k$ .
- The autocorrelations are the correlations associated with these scatterplots.

# Autocorrelation and Autocovariance

- Autocovariance at lag  $k$

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

- Autocorrelation at lag  $k$

$$r_k = c_k / c_0$$

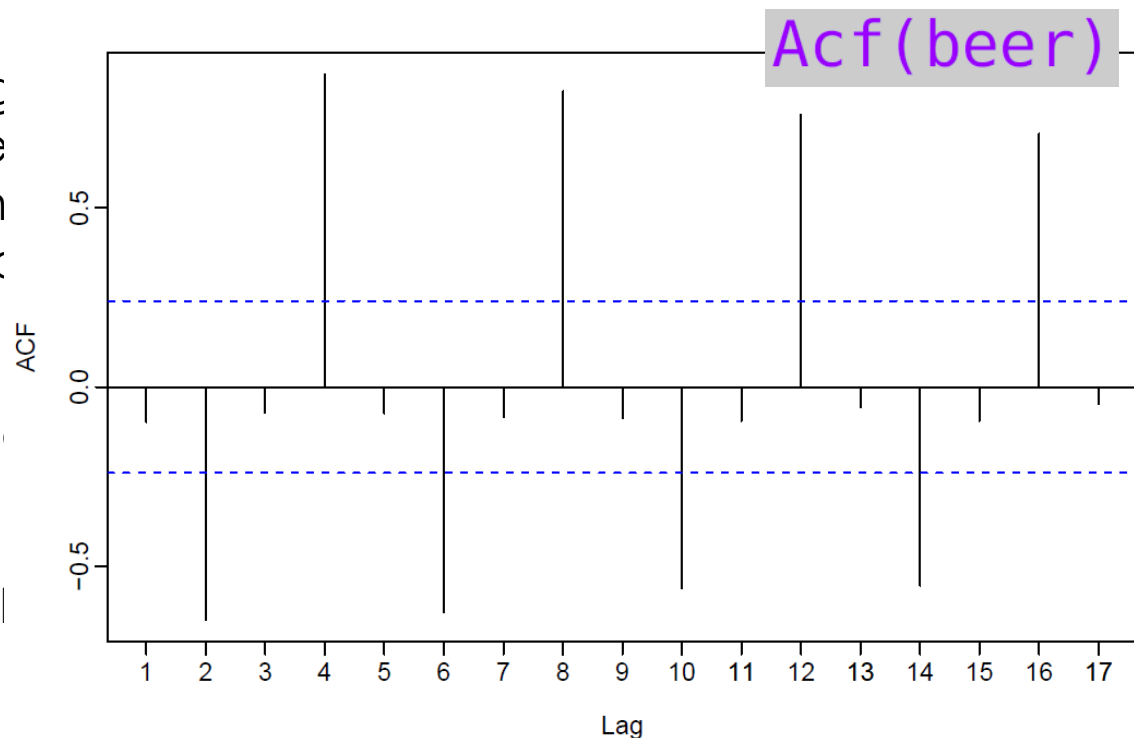
- $r_1$  indicates how successive values of  $y$  relate to each other
- $r_2$  indicates how  $y$  values two periods apart relate to each other
- $r_k$  is like the correlation between  $y_t$  and  $y_{t-k}$

# Example: Autocorrelation & Correlogram

Results for first 9 lags for beer data:

$r_1$	$r_2$	$r_3$	$r_4$	$r_5$	$r_6$	$r_7$	$r_8$	$r_9$
-0.126	-0.650	-0.094	0.863	-0.099	-0.642	-0.098	0.834	-0.116

- $r_4$  higher than for the other lag
  - due to the seasonal pattern in the data: the peaks tend to be 4 quarters apart and the troughs tend to be 2 quarters apart.
- $r_2$  is more negative than for the other lags
  - because troughs tend to be 2 quarters behind peaks



# Detecting Trend & Seasonality via ACF

---

- If there is seasonality, the **ACF at the seasonal lag** (e.g., 12 for monthly data) will be **large** and **positive**:
  - For seasonal **monthly** data, a large ACF value will be seen at lag 12 and possibly also at lags 24, 36, . . .
  - For seasonal **quarterly** data, a large ACF value will be seen at lag 4 and possibly also at lags 8, 12, . . .
- The slowly decaying ACF indicates trend


# Autocorrelation with Acf

- In R forecast package Acf function is used to compute and plot autocorrelations of a series at different lags

```
library("forecast")
library("zoo")

Amtrak.data <- read.csv("Amtrak data.csv")
ridership.ts <- ts(Amtrak.data$Ridership, start = c(1991, 1), end = c(2004, 3), freq = 12)
ridership.24.ts <- window(ridership.ts, start = c(1991, 1), end = c(1991, 24))

Acf(ridership.24.ts, lag.max = 12, main = "")
```






# Autocorrelation Interpretation

---


- Strong autocorrelation (positive or negative) at a lag larger than 1 typically reflects a cyclical pattern.
  - strong positive autocorrelation at lag-12 in monthly data reflects an annual seasonality where values during a given month each year are positively correlated
- Positive lag-1 autocorrelation (called “stickiness”) describes a series where consecutive values move generally in the same direction.
  - In the presence of a strong linear trend, we would expect to see a strong and positive lag-1 autocorrelation
- Negative lag-1 autocorrelation reflects swings in the series, where high values are immediately followed by low values and vice versa

# Autocorrelation detection of seasonality patterns

- The strongest negative correlation at lag-6 indicates a biannual pattern in ridership, with 6-month switches from high to low
  - high summer and low-winter pattern




# Ex: What ACF says about *Monthly Electricity Production in Australia*?




# What ACF tells us about TS data sets?


1. Daily morning temperature of a cow




2. Accidental deaths in USA (monthly)




3. International airline passengers




4. Annual mink trappings (Canada)




A




B



C



D



# Autocorrelation: Residual Series

- It is useful to look at autocorrelations of residual series.
  - e.g. after fitting a regression model we can examine the autocorrelation of the series of residuals
- If we adequately modeled the seasonal pattern, then the residual series should show no autocorrelation at the season's lag.
  - 6 month and 12 month cyclical behavior is not there in the residual series from the regression with seasonality and quadratic trend

